

Name: Kunal Singhal University Roll No. : 2315510107 Section: CB

**Assignment Title:** End-to-End ANN Model Implementation and Analysis Instructions:

1. Follow the tasks step by step in this notebook.
2. Write your code in the provided cells only.
3. Submit the following:

Completed .ipynb file

**A report (.pdf or .docx) with answers to the analysis questions provided at the end.**

Task 1: Import Required Libraries

In [ ]:

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
```

In [44]:

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

Task 2: Load Dataset (Use from sklearn.datasets or any open dataset)

In [86]:

```
df = pd.read_csv("Dataset.csv")
```

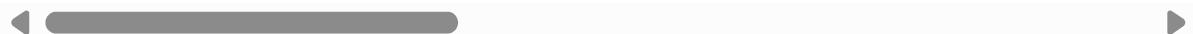
In [87]:

```
df.head()
```

Out[87]:

	Surname	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMe
0	2023	668.0	33.0	3.0	0.00	2.0	1	
1	2024	627.0	33.0	1.0	0.00	2.0	1	
2	1236	678.0	40.0	10.0	0.00	2.0	1	
3	1362	581.0	34.0	2.0	148882.54	1.0	1	
4	491	716.0	33.0	5.0	0.00	2.0	1	

5 rows × 25 columns



In [88]: `df.tail()`

Out[88]:

	Surname	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsAct
175023	2226	644.0	28.0	7.0	155060.41	1.0	1	
175024	2846	800.0	29.0	2.0	0.00	2.0	0	
175025	1999	771.0	39.0	5.0	0.00	2.0	1	
175026	1336	516.0	35.0	10.0	57369.61	1.0	1	
175027	1570	709.0	36.0	7.0	0.00	1.0	0	

5 rows × 25 columns



### ✓ Task 3: Perform Exploratory Data Analysis (EDA)

1. Show basic info (shape, columns, datatypes)
2. Describe data
3. Use visualizations (histograms, boxplots, heatmaps, pairplots etc.)

Identify class imbalance if any

In [90]: `df.shape`

Out[90]: (175028, 25)

In [91]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 175028 entries, 0 to 175027
Data columns (total 25 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Surname          175028 non-null   int64  
 1   CreditScore      175028 non-null   float64 
 2   Age              175028 non-null   float64 
 3   Tenure           175028 non-null   float64 
 4   Balance          175028 non-null   float64 
 5   NumOfProducts    175028 non-null   float64 
 6   HasCrCard        175028 non-null   int64  
 7   IsActiveMember   175028 non-null   int64  
 8   EstimatedSalary  175028 non-null   float64 
 9   Exited           175028 non-null   int64  
 10  Surname_tfidf_0  175028 non-null   float64 
 11  Surname_tfidf_1  175028 non-null   float64 
 12  Surname_tfidf_2  175028 non-null   float64 
 13  Surname_tfidf_3  175028 non-null   float64 
 14  Surname_tfidf_4  175028 non-null   float64 
 15  France           175028 non-null   int64  
 16  Germany          175028 non-null   int64  
 17  Spain            175028 non-null   int64  
 18  Female           175028 non-null   int64  
 19  Male              175028 non-null   int64  
 20  Mem_no_Products  175028 non-null   float64 
 21  Cred_Bal_Sal    175028 non-null   float64 
 22  Bal_sal          175028 non-null   float64 
 23  Tenure_Age       175028 non-null   float64 
 24  Age_Tenure_product 175028 non-null   float64 
dtypes: float64(16), int64(9)
memory usage: 33.4 MB

```

In [92]: `df.describe()`

	Surname	CreditScore	Age	Tenure	Balance	NumOf
<b>count</b>	175028.000000	175028.000000	175028.000000	175028.000000	175028.000000	17502
<b>mean</b>	1617.462812	656.113068	38.171337	5.019968	56676.772103	
<b>std</b>	813.852174	81.148273	8.969565	2.811155	62982.418525	
<b>min</b>	0.000000	350.000000	18.000000	0.000000	0.000000	
<b>25%</b>	949.000000	597.000000	32.000000	3.000000	0.000000	
<b>50%</b>	1666.000000	659.000000	37.000000	5.000000	0.000000	
<b>75%</b>	2292.000000	710.000000	42.000000	7.000000	120727.970000	
<b>max</b>	2931.000000	850.000000	92.000000	10.000000	250898.090000	

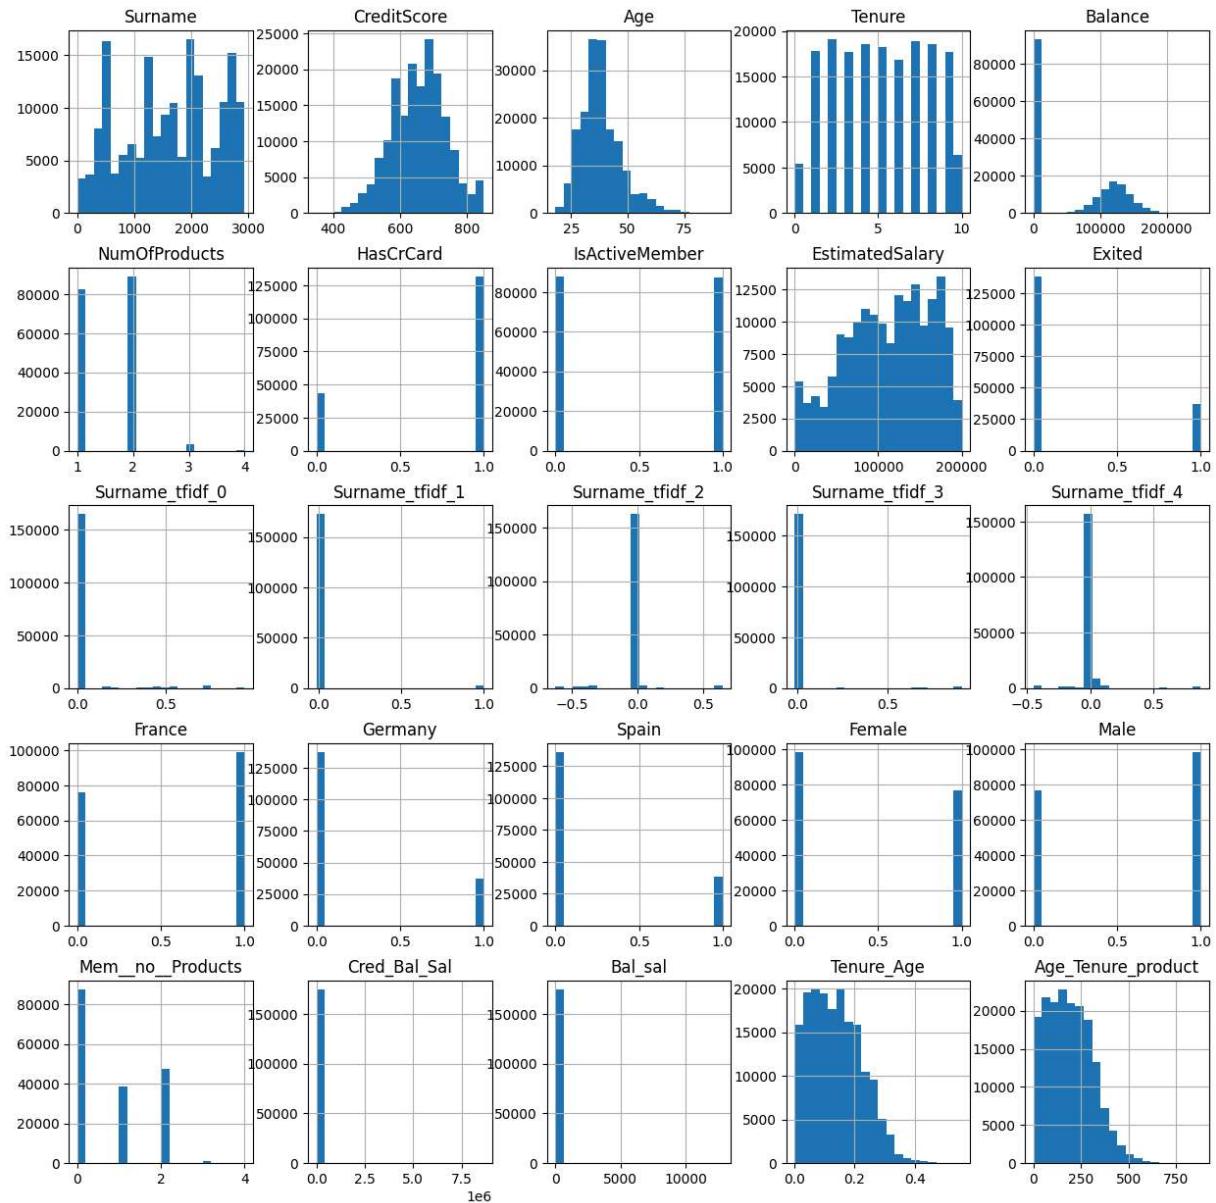
8 rows × 25 columns



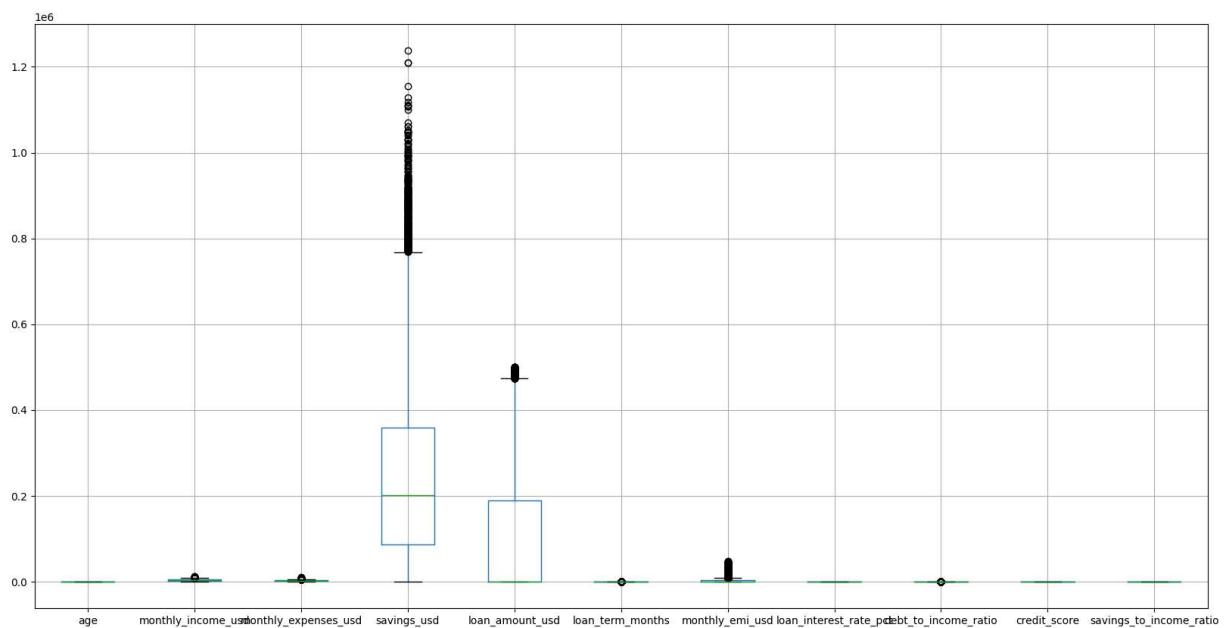
In [93]: `df.columns`

Out[93]: Index(['Surname', 'CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited', 'Surname\_tfidf\_0', 'Surname\_tfidf\_1', 'Surname\_tfidf\_2', 'Surname\_tfidf\_3', 'Surname\_tfidf\_4', 'France', 'Germany', 'Spain', 'Female', 'Male', 'Mem\_no\_Products', 'Cred\_Bal\_Sal', 'Bal\_sal', 'Tenure\_Age', 'Age\_Tenure\_product'],  
 dtype='object')

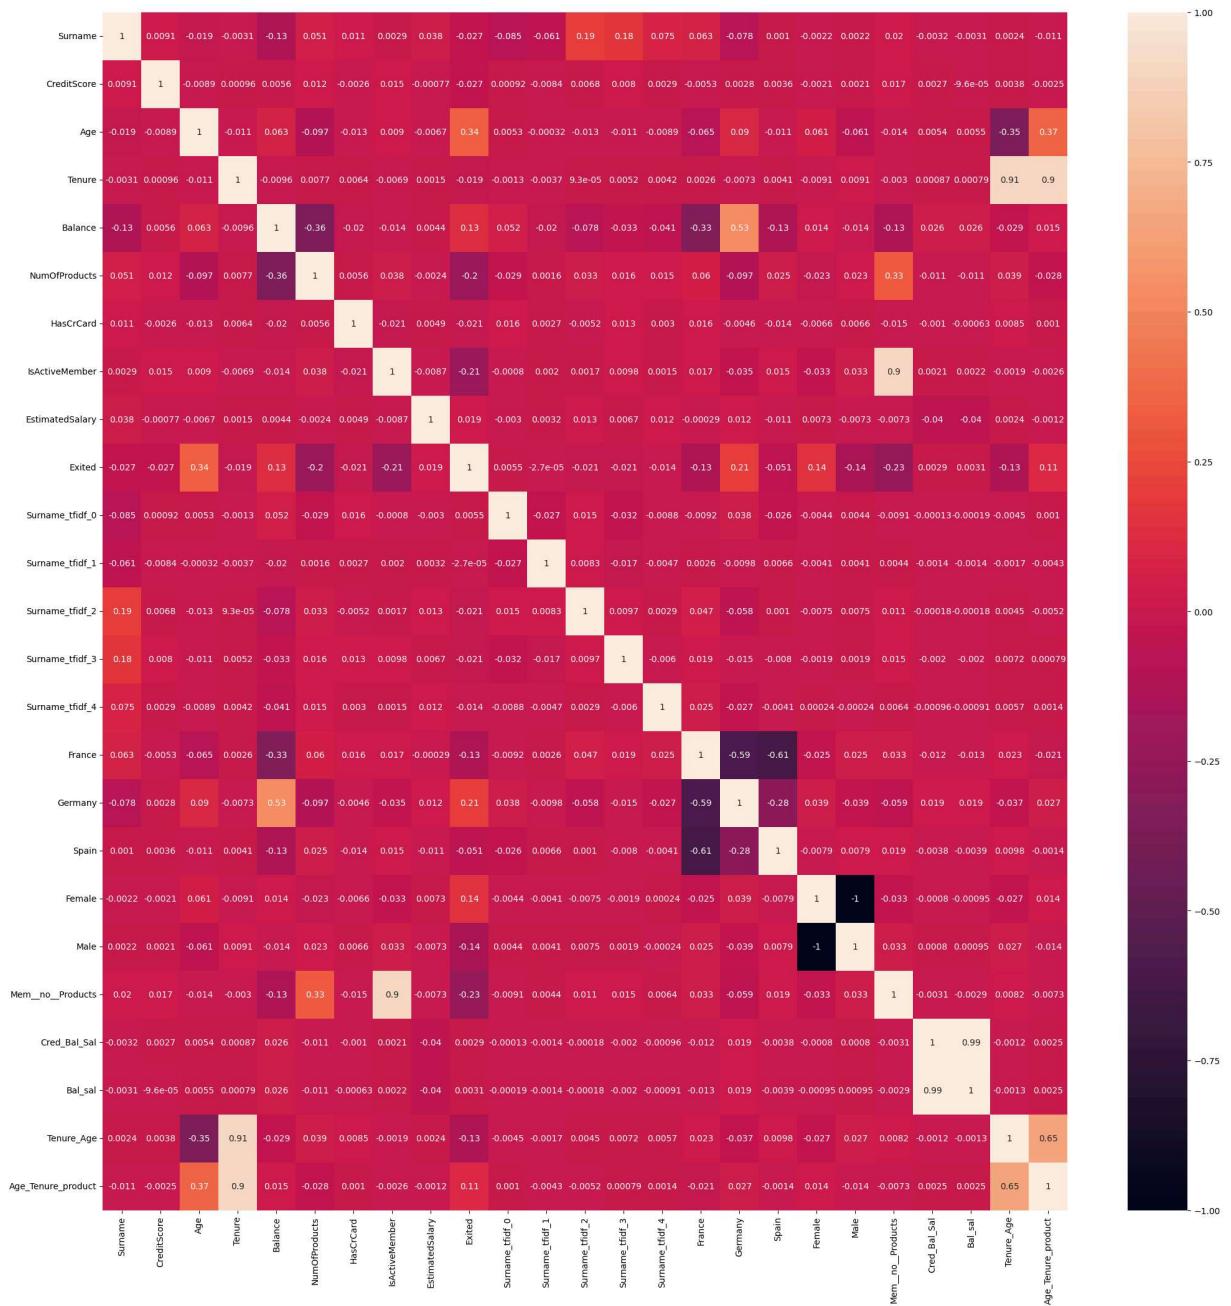
In [95]: `df.hist(bins=20, figsize=(15,15))  
 plt.show()`



In [54]: `df.boxplot(figsize=(20,10))  
 plt.show()`



```
In [98]: corr_matrix = df.corr(numeric_only=True)
plt.figure(figsize=(25, 25))
sns.heatmap(corr_matrix, annot=True, )
plt.show()
```



#### ✓ Task 4: Feature Engineering

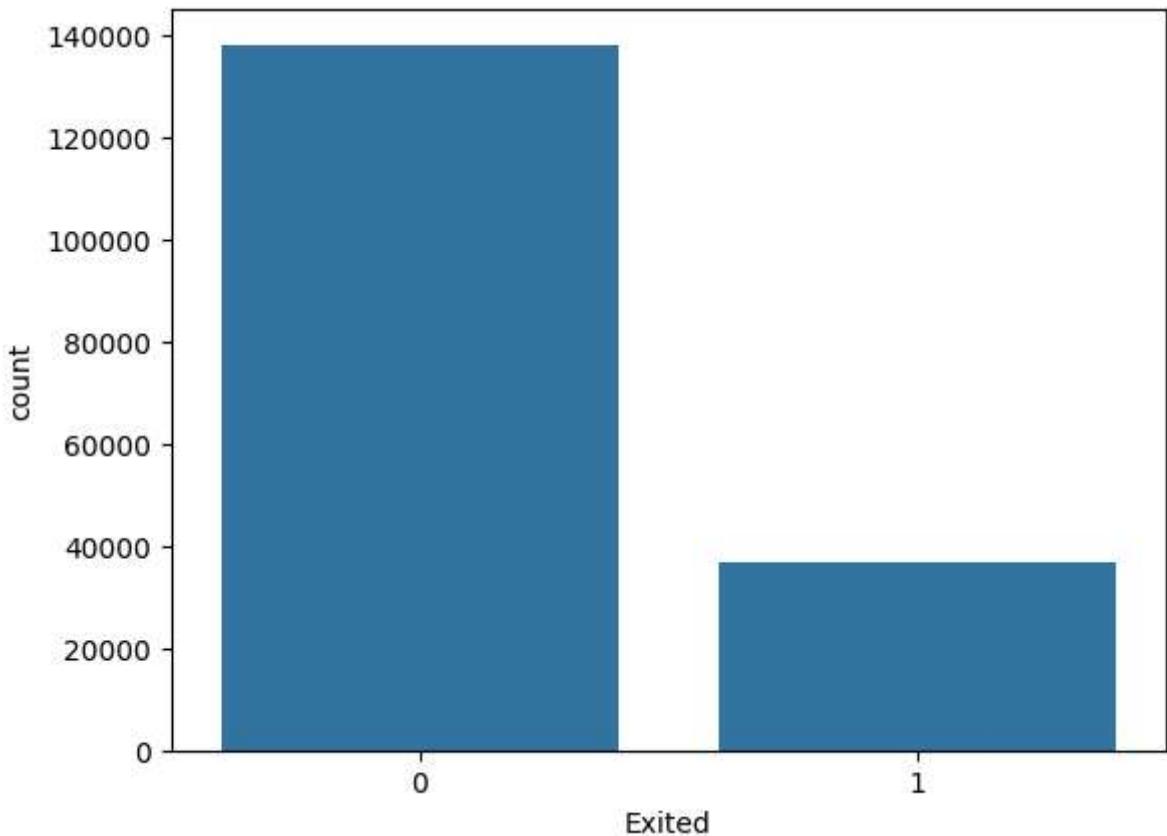
1. Handle missing values if any
2. Encode categorical variables if any
3. Create or remove irrelevant features

In [99]: `df.isnull().sum()`

```
Out[99]: Surname      0  
CreditScore      0  
Age             0  
Tenure          0  
Balance         0  
NumOfProducts    0  
HasCrCard       0  
IsActiveMember   0  
EstimatedSalary  0  
Exited          0  
Surname_tfidf_0  0  
Surname_tfidf_1  0  
Surname_tfidf_2  0  
Surname_tfidf_3  0  
Surname_tfidf_4  0  
France          0  
Germany         0  
Spain           0  
Female          0  
Male            0  
Mem_no_Products 0  
Cred_Bal_Sal    0  
Bal_sal         0  
Tenure_Age      0  
Age_Tenure_product 0  
dtype: int64
```

```
In [101...]: sns.countplot(x='Exited', data=df)  
df['Exited'].value_counts()
```

```
Out[101...]: Exited  
0    138071  
1    36957  
Name: count, dtype: int64
```



Task 5: Feature Scaling / Normalization

1. StandardScaler / MinMaxScaler

Task 6: Split Data into Training and Testing Sets (80/20 or 70/30)

```
In [103]: X = df.drop(['Exited'], axis=1)  
y = df['Exited']
```

```
In [105]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [62]: X_test.shape
```

```
Out[62]: (9728, 32)
```

```
In [107]: sc = StandardScaler()  
X_train = sc.fit_transform(X_train)  
X_test = sc.transform(X_test)
```

Task 7: Define the ANN Model (Sequential API - Keras)

1. Design architecture (input, hidden, output layers)
2. Use appropriate activation functions

```
In [109]: model = Sequential([  
    Dense(128, activation='relu'),  
    Dense(64, activation='relu'),
```

```
Dense(32,activation='relu'),  
Dense(1,activation='sigmoid')  
])
```

In [110... model.summary()

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
dense_20 (Dense)	?	0 (unbuilt)
dense_21 (Dense)	?	0 (unbuilt)
dense_22 (Dense)	?	0 (unbuilt)
dense_23 (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)

Task 8: Compile the Model

- Select optimizer, loss function, and evaluation metrics

In [115... model.compile(optimizer='adam', loss='binary\_crossentropy', metrics=['accuracy'])

Task 9: Train the Model

- Use model.fit()
- Track performance using training history

In [112... m = model.fit(X\_train, y\_train, epochs=50, validation\_split=0.2)

Epoch 1/50  
**3501/3501** 10s 2ms/step - accuracy: 0.8551 - loss: 0.3436 - val\_accuracy: 0.8617 - val\_loss: 0.3324  
Epoch 2/50  
**3501/3501** 14s 4ms/step - accuracy: 0.8624 - loss: 0.3282 - val\_accuracy: 0.8619 - val\_loss: 0.3333  
Epoch 3/50  
**3501/3501** 14s 4ms/step - accuracy: 0.8637 - loss: 0.3264 - val\_accuracy: 0.8605 - val\_loss: 0.3307  
Epoch 4/50  
**3501/3501** 14s 4ms/step - accuracy: 0.8644 - loss: 0.3250 - val\_accuracy: 0.8620 - val\_loss: 0.3309  
Epoch 5/50  
**3501/3501** 14s 4ms/step - accuracy: 0.8647 - loss: 0.3238 - val\_accuracy: 0.8621 - val\_loss: 0.3304  
Epoch 6/50  
**3501/3501** 7s 2ms/step - accuracy: 0.8646 - loss: 0.3232 - val\_accuracy: 0.8633 - val\_loss: 0.3279  
Epoch 7/50  
**3501/3501** 7s 2ms/step - accuracy: 0.8650 - loss: 0.3222 - val\_accuracy: 0.8638 - val\_loss: 0.3276  
Epoch 8/50  
**3501/3501** 13s 4ms/step - accuracy: 0.8645 - loss: 0.3216 - val\_accuracy: 0.8619 - val\_loss: 0.3304  
Epoch 9/50  
**3501/3501** 14s 4ms/step - accuracy: 0.8652 - loss: 0.3207 - val\_accuracy: 0.8629 - val\_loss: 0.3280  
Epoch 10/50  
**3501/3501** 14s 4ms/step - accuracy: 0.8658 - loss: 0.3201 - val\_accuracy: 0.8630 - val\_loss: 0.3287  
Epoch 11/50  
**3501/3501** 14s 4ms/step - accuracy: 0.8657 - loss: 0.3195 - val\_accuracy: 0.8625 - val\_loss: 0.3277  
Epoch 12/50  
**3501/3501** 7s 2ms/step - accuracy: 0.8661 - loss: 0.3191 - val\_accuracy: 0.8634 - val\_loss: 0.3286  
Epoch 13/50  
**3501/3501** 13s 4ms/step - accuracy: 0.8664 - loss: 0.3183 - val\_accuracy: 0.8621 - val\_loss: 0.3284  
Epoch 14/50  
**3501/3501** 16s 5ms/step - accuracy: 0.8666 - loss: 0.3176 - val\_accuracy: 0.8636 - val\_loss: 0.3296  
Epoch 15/50  
**3501/3501** 7s 2ms/step - accuracy: 0.8668 - loss: 0.3169 - val\_accuracy: 0.8641 - val\_loss: 0.3279  
Epoch 16/50  
**3501/3501** 6s 2ms/step - accuracy: 0.8675 - loss: 0.3164 - val\_accuracy: 0.8640 - val\_loss: 0.3281  
Epoch 17/50  
**3501/3501** 6s 2ms/step - accuracy: 0.8671 - loss: 0.3162 - val\_accuracy: 0.8601 - val\_loss: 0.3317  
Epoch 18/50  
**3501/3501** 6s 2ms/step - accuracy: 0.8674 - loss: 0.3149 - val\_accuracy: 0.8622 - val\_loss: 0.3326  
Epoch 19/50  
**3501/3501** 6s 2ms/step - accuracy: 0.8678 - loss: 0.3147 - val\_accuracy: 0.8623 - val\_loss: 0.3328

```
accuracy: 0.8622 - val_loss: 0.3305
Epoch 20/50
3501/3501 10s 1ms/step - accuracy: 0.8684 - loss: 0.3136 - val_
accuracy: 0.8621 - val_loss: 0.3308
Epoch 21/50
3501/3501 6s 2ms/step - accuracy: 0.8677 - loss: 0.3133 - val_a
ccuracy: 0.8626 - val_loss: 0.3298
Epoch 22/50
3501/3501 7s 2ms/step - accuracy: 0.8680 - loss: 0.3122 - val_a
ccuracy: 0.8611 - val_loss: 0.3328
Epoch 23/50
3501/3501 6s 2ms/step - accuracy: 0.8690 - loss: 0.3118 - val_a
ccuracy: 0.8618 - val_loss: 0.3320
Epoch 24/50
3501/3501 6s 2ms/step - accuracy: 0.8690 - loss: 0.3115 - val_a
ccuracy: 0.8617 - val_loss: 0.3338
Epoch 25/50
3501/3501 6s 2ms/step - accuracy: 0.8690 - loss: 0.3103 - val_a
ccuracy: 0.8610 - val_loss: 0.3336
Epoch 26/50
3501/3501 7s 2ms/step - accuracy: 0.8696 - loss: 0.3098 - val_a
ccuracy: 0.8613 - val_loss: 0.3346
Epoch 27/50
3501/3501 7s 2ms/step - accuracy: 0.8701 - loss: 0.3089 - val_a
ccuracy: 0.8631 - val_loss: 0.3357
Epoch 28/50
3501/3501 10s 2ms/step - accuracy: 0.8695 - loss: 0.3083 - val_
accuracy: 0.8617 - val_loss: 0.3376
Epoch 29/50
3501/3501 7s 2ms/step - accuracy: 0.8700 - loss: 0.3077 - val_a
ccuracy: 0.8611 - val_loss: 0.3353
Epoch 30/50
3501/3501 6s 2ms/step - accuracy: 0.8707 - loss: 0.3069 - val_a
ccuracy: 0.8612 - val_loss: 0.3357
Epoch 31/50
3501/3501 6s 2ms/step - accuracy: 0.8714 - loss: 0.3063 - val_a
ccuracy: 0.8591 - val_loss: 0.3434
Epoch 32/50
3501/3501 6s 2ms/step - accuracy: 0.8706 - loss: 0.3056 - val_a
ccuracy: 0.8592 - val_loss: 0.3379
Epoch 33/50
3501/3501 14s 4ms/step - accuracy: 0.8720 - loss: 0.3051 - val_
accuracy: 0.8592 - val_loss: 0.3393
Epoch 34/50
3501/3501 14s 4ms/step - accuracy: 0.8719 - loss: 0.3042 - val_
accuracy: 0.8593 - val_loss: 0.3393
Epoch 35/50
3501/3501 14s 4ms/step - accuracy: 0.8717 - loss: 0.3035 - val_
accuracy: 0.8610 - val_loss: 0.3395
Epoch 36/50
3501/3501 15s 4ms/step - accuracy: 0.8721 - loss: 0.3029 - val_
accuracy: 0.8607 - val_loss: 0.3397
Epoch 37/50
3501/3501 15s 4ms/step - accuracy: 0.8722 - loss: 0.3024 - val_
accuracy: 0.8578 - val_loss: 0.3432
Epoch 38/50
```

```
3501/3501 ━━━━━━━━ 12s 4ms/step - accuracy: 0.8725 - loss: 0.3015 - val_
accuracy: 0.8600 - val_loss: 0.3441
Epoch 39/50
3501/3501 ━━━━━━━━ 14s 2ms/step - accuracy: 0.8727 - loss: 0.3011 - val_
accuracy: 0.8577 - val_loss: 0.3434
Epoch 40/50
3501/3501 ━━━━━━ 6s 2ms/step - accuracy: 0.8734 - loss: 0.3006 - val_a
ccuracy: 0.8604 - val_loss: 0.3443
Epoch 41/50
3501/3501 ━━━━━━ 6s 2ms/step - accuracy: 0.8733 - loss: 0.2996 - val_a
ccuracy: 0.8598 - val_loss: 0.3449
Epoch 42/50
3501/3501 ━━━━━━ 6s 2ms/step - accuracy: 0.8737 - loss: 0.2988 - val_a
ccuracy: 0.8577 - val_loss: 0.3465
Epoch 43/50
3501/3501 ━━━━━━ 6s 2ms/step - accuracy: 0.8744 - loss: 0.2983 - val_a
ccuracy: 0.8577 - val_loss: 0.3492
Epoch 44/50
3501/3501 ━━━━━━ 6s 2ms/step - accuracy: 0.8738 - loss: 0.2978 - val_a
ccuracy: 0.8573 - val_loss: 0.3451
Epoch 45/50
3501/3501 ━━━━━━ 6s 2ms/step - accuracy: 0.8737 - loss: 0.2975 - val_a
ccuracy: 0.8570 - val_loss: 0.3514
Epoch 46/50
3501/3501 ━━━━━━ 6s 2ms/step - accuracy: 0.8745 - loss: 0.2968 - val_a
ccuracy: 0.8590 - val_loss: 0.3482
Epoch 47/50
3501/3501 ━━━━━━ 6s 2ms/step - accuracy: 0.8747 - loss: 0.2960 - val_a
ccuracy: 0.8590 - val_loss: 0.3517
Epoch 48/50
3501/3501 ━━━━━━ 6s 2ms/step - accuracy: 0.8750 - loss: 0.2968 - val_a
ccuracy: 0.8583 - val_loss: 0.3498
Epoch 49/50
3501/3501 ━━━━━━ 6s 2ms/step - accuracy: 0.8758 - loss: 0.2948 - val_a
ccuracy: 0.8562 - val_loss: 0.3538
Epoch 50/50
3501/3501 ━━━━━━ 6s 2ms/step - accuracy: 0.8760 - loss: 0.2944 - val_a
ccuracy: 0.8570 - val_loss: 0.3539
```

In [114...]

`model.summary()`

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
dense_20 (Dense)	(None, 128)	3,200
dense_21 (Dense)	(None, 64)	8,256
dense_22 (Dense)	(None, 32)	2,080
dense_23 (Dense)	(None, 1)	33

Total params: 40,709 (159.02 KB)

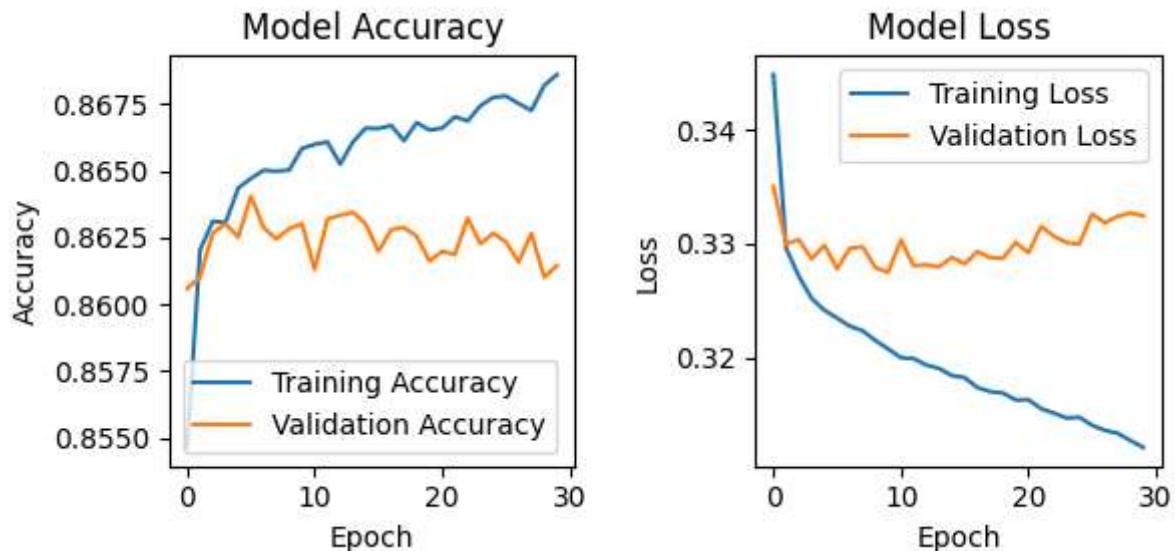
Trainable params: 13,569 (53.00 KB)

**Non-trainable params:** 0 (0.00 B)  
**Optimizer params:** 27,140 (106.02 KB)

Task 10: Visualize Model Performance

- Plot loss and accuracy curves for training and validation

```
In [138...]: plt.figure(figsize=(6,3))
plt.subplot(1, 2, 1)
plt.plot(m.history['accuracy'], label='Training Accuracy')
plt.plot(m.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.subplot(1, 2, 2)
plt.plot(m.history['loss'], label='Training Loss')
plt.plot(m.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()
plt.show()
```



In [ ]:

Task 11: Evaluate Model on Test Data

- Accuracy, loss
- Predict test set

```
In [119...]: loss, mae = model.evaluate(X_test,y_test)
```

1094/1094 ━━━━━━━━ 2s 1ms/step - accuracy: 0.8591 - loss: 0.3481

```
In [120...]: y_pred = model.predict(X_test)
```

1094/1094 ━━━━━━ 1s 917us/step

✓ Task 12: Confusion Matrix and Classification Report

- Use sklearn metrics to analyze performance

```
In [121...]: mse = mean_squared_error(y_test, y_pred)
mse
```

Out[121...]: 0.10393407940864563

```
In [123...]: mae = mean_absolute_error(y_test, y_pred)
mae
```

Out[123...]: 0.1914944052696228

```
In [125...]: rmse = np.sqrt(mse)
rmse
```

Out[125...]: np.float64(0.32238808819285747)

```
In [126...]: r2 = r2_score(y_test, y_pred)
r2
```

Out[126...]: 0.3765411376953125

✓ Task 13: Hyperparameter Tuning (Try any of the following)

- Vary batch size, epochs, optimizers
- Use Keras Tuner or manual tuning

```
In [131...]: model = Sequential([
    Dense(128, activation='relu', input_shape=(X_train.shape[1],)),
    Dense(32, activation='relu'),
    Dense(1, activation='sigmoid')
])
```

c:\Users\This PC\Desktop\DeepLearning\.venv\Lib\site-packages\keras\src\layers\core\dense.py:92: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.  
super().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)

```
In [132...]: model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model.summary()
```

**Model: "sequential\_7"**

Layer (type)	Output Shape	Param #
dense_28 (Dense)	(None, 128)	3,200
dense_29 (Dense)	(None, 32)	4,128
dense_30 (Dense)	(None, 1)	33

Total params: 7,361 (28.75 KB)

Trainable params: 7,361 (28.75 KB)

Non-trainable params: 0 (0.00 B)

```
In [133]: m = model.fit(X_train, y_train, epochs=30, validation_split=0.2)
```

Epoch 1/30

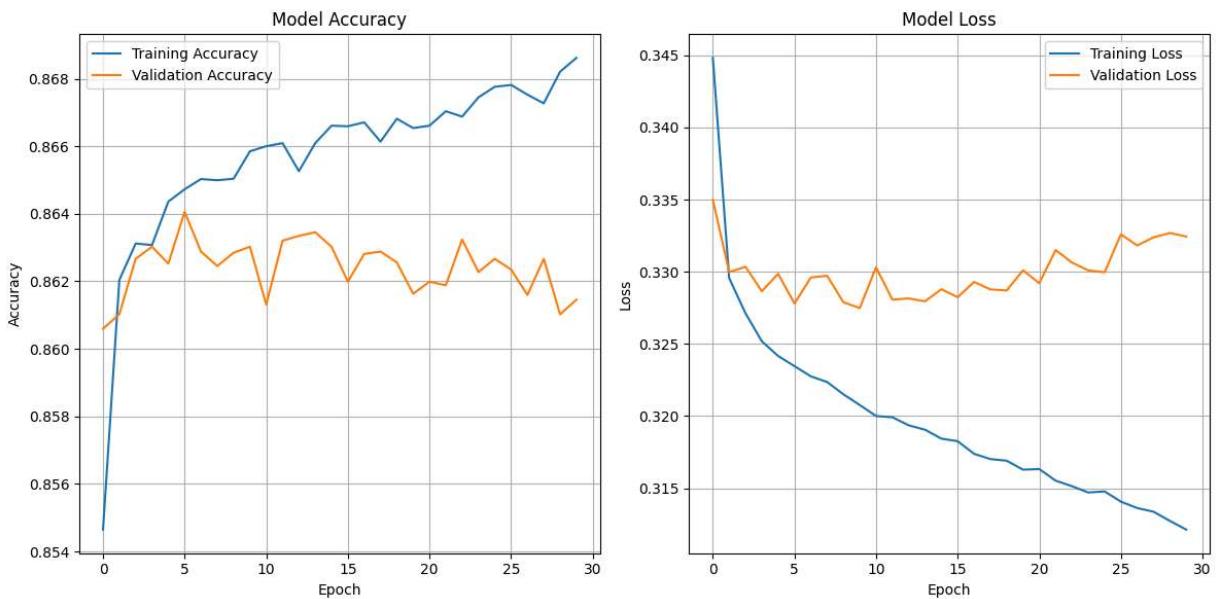
```
3501/3501 ━━━━━━━━ 7s 2ms/step - accuracy: 0.8546 - loss: 0.3448 - val_accuracy: 0.8606 - val_loss: 0.3350
Epoch 2/30
3501/3501 ━━━━━━━━ 6s 2ms/step - accuracy: 0.8620 - loss: 0.3296 - val_accuracy: 0.8610 - val_loss: 0.3300
Epoch 3/30
3501/3501 ━━━━━━━━ 5s 2ms/step - accuracy: 0.8631 - loss: 0.3271 - val_accuracy: 0.8627 - val_loss: 0.3303
Epoch 4/30
3501/3501 ━━━━━━━━ 5s 2ms/step - accuracy: 0.8631 - loss: 0.3252 - val_accuracy: 0.8630 - val_loss: 0.3287
Epoch 5/30
3501/3501 ━━━━━━━━ 5s 1ms/step - accuracy: 0.8644 - loss: 0.3242 - val_accuracy: 0.8625 - val_loss: 0.3299
Epoch 6/30
3501/3501 ━━━━━━━━ 5s 2ms/step - accuracy: 0.8647 - loss: 0.3235 - val_accuracy: 0.8641 - val_loss: 0.3278
Epoch 7/30
3501/3501 ━━━━━━━━ 5s 1ms/step - accuracy: 0.8650 - loss: 0.3228 - val_accuracy: 0.8629 - val_loss: 0.3296
Epoch 8/30
3501/3501 ━━━━━━━━ 5s 1ms/step - accuracy: 0.8650 - loss: 0.3224 - val_accuracy: 0.8625 - val_loss: 0.3297
Epoch 9/30
3501/3501 ━━━━━━━━ 5s 1ms/step - accuracy: 0.8650 - loss: 0.3215 - val_accuracy: 0.8628 - val_loss: 0.3279
Epoch 10/30
3501/3501 ━━━━━━━━ 6s 2ms/step - accuracy: 0.8659 - loss: 0.3208 - val_accuracy: 0.8630 - val_loss: 0.3275
Epoch 11/30
3501/3501 ━━━━━━━━ 5s 2ms/step - accuracy: 0.8660 - loss: 0.3200 - val_accuracy: 0.8613 - val_loss: 0.3303
Epoch 12/30
3501/3501 ━━━━━━━━ 5s 2ms/step - accuracy: 0.8661 - loss: 0.3199 - val_accuracy: 0.8632 - val_loss: 0.3281
Epoch 13/30
3501/3501 ━━━━━━━━ 5s 1ms/step - accuracy: 0.8653 - loss: 0.3194 - val_accuracy: 0.8633 - val_loss: 0.3282
Epoch 14/30
3501/3501 ━━━━━━━━ 6s 2ms/step - accuracy: 0.8661 - loss: 0.3191 - val_accuracy: 0.8635 - val_loss: 0.3280
Epoch 15/30
3501/3501 ━━━━━━━━ 6s 2ms/step - accuracy: 0.8666 - loss: 0.3184 - val_accuracy: 0.8630 - val_loss: 0.3288
Epoch 16/30
3501/3501 ━━━━━━━━ 5s 1ms/step - accuracy: 0.8666 - loss: 0.3183 - val_accuracy: 0.8620 - val_loss: 0.3282
Epoch 17/30
3501/3501 ━━━━━━━━ 6s 2ms/step - accuracy: 0.8667 - loss: 0.3174 - val_accuracy: 0.8628 - val_loss: 0.3293
Epoch 18/30
3501/3501 ━━━━━━━━ 5s 2ms/step - accuracy: 0.8661 - loss: 0.3170 - val_accuracy: 0.8629 - val_loss: 0.3288
Epoch 19/30
3501/3501 ━━━━━━━━ 7s 2ms/step - accuracy: 0.8668 - loss: 0.3169 - val_accuracy: 0.8626 - val_loss: 0.3287
```

```
Epoch 20/30
3501/3501 5s 2ms/step - accuracy: 0.8665 - loss: 0.3163 - val_accuracy: 0.8616 - val_loss: 0.3301
Epoch 21/30
3501/3501 6s 2ms/step - accuracy: 0.8666 - loss: 0.3163 - val_accuracy: 0.8620 - val_loss: 0.3292
Epoch 22/30
3501/3501 6s 2ms/step - accuracy: 0.8670 - loss: 0.3155 - val_accuracy: 0.8619 - val_loss: 0.3315
Epoch 23/30
3501/3501 6s 2ms/step - accuracy: 0.8669 - loss: 0.3151 - val_accuracy: 0.8632 - val_loss: 0.3306
Epoch 24/30
3501/3501 6s 2ms/step - accuracy: 0.8674 - loss: 0.3147 - val_accuracy: 0.8623 - val_loss: 0.3301
Epoch 25/30
3501/3501 5s 1ms/step - accuracy: 0.8678 - loss: 0.3148 - val_accuracy: 0.8627 - val_loss: 0.3300
Epoch 26/30
3501/3501 5s 1ms/step - accuracy: 0.8678 - loss: 0.3141 - val_accuracy: 0.8623 - val_loss: 0.3326
Epoch 27/30
3501/3501 5s 1ms/step - accuracy: 0.8675 - loss: 0.3136 - val_accuracy: 0.8616 - val_loss: 0.3318
Epoch 28/30
3501/3501 5s 2ms/step - accuracy: 0.8673 - loss: 0.3134 - val_accuracy: 0.8627 - val_loss: 0.3324
Epoch 29/30
3501/3501 5s 1ms/step - accuracy: 0.8682 - loss: 0.3127 - val_accuracy: 0.8610 - val_loss: 0.3327
Epoch 30/30
3501/3501 5s 1ms/step - accuracy: 0.8686 - loss: 0.3121 - val_accuracy: 0.8615 - val_loss: 0.3324
```

In [134...]

```
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(m.history['accuracy'], label='Training Accuracy')
plt.plot(m.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(m.history['loss'], label='Training Loss')
plt.plot(m.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



Task 14: Save and Load the Model (Optional)

Report Questions (To be submitted separately)

- Dataset Selection: What dataset did you choose? Describe its purpose and target.
- EDA Findings: What key insights did you gain during EDA?
- Feature Engineering: What transformations did you apply? Why?
- ANN Architecture: List number of layers, neurons, activation functions.
- Performance Summary: What were the final accuracy and loss?
- Confusion Matrix Analysis: What did it reveal about model behavior?
- Hyperparameter Tuning: What parameters did you change and what was the result?
- Improvement Ideas: Suggest one way to improve performance.
- Overfitting/Underfitting: Did you face it? How did you address it?